

Introduction to Bayesian statistics with R

5. Case studies and GLMMs

Olivier Gimenez

last updated: 2025-03-11

What is a Generalized Linear Model (or GLM)?

Survival of passengers on the Titanic ~ Class

Read titanic_long.csv dataset.

```
titanic <- read.csv("dat/titanic_long.csv")
```

```
head(titanic)
```

```
#>   class  age  sex survived
```

```
#> 1 first adult male        1
```

```
#> 2 first adult male        1
```

```
#> 3 first adult male        1
```

```
#> 4 first adult male        1
```


```
#> 5 first adult male        1
```

```
#> 6 first adult male        1
```

Let's fit a linear model

```
titanic.lm <- lm(survived ~ class, data = titanic)
```

Clearly, the residuals are not normal!



Generalized linear models (GLMs)

- GLMs extend the linear model to scenarios that involve **non-normal error distributions**, hence the term **generalized**
- The **mean response** is expressed as a **linear function of the predictors** using a **link function**, hence the term **linear**

Generalized Linear Models

1. Response variable

- Bernoulli/Binomial: Binary variables 0/1
- Poisson: Counts 0, 1, 2, ...
- Normal: Real values
- etc

2. Predictors (continuous or categorical)

3. Link function

- Gaussian: identity
- Binomial: logit
- Poisson: log
- Type in ?family

Bernoulli/Binomial distribution (logistic regression)

- Response variable: Yes/No (e.g. dead/alive, male/female, presence/absence)
- Link function: `logit`

$$\text{logit}(p) = \ln \left(\frac{p}{1-p} \right)$$

- Then, if predictor is x

Response \sim Distribution(Mean Response)

$$Y_i \sim \text{Bernoulli}(p_i)$$

$$\text{logit}(p_i) = a + b x_i$$

$$p_i = \text{logit}^{-1}(a + b x_i) = \frac{e^{a+b x_i}}{1 + e^{a+b x_i}}$$

Poisson distributoin (Poisson regression)

- Response variable: Counts (0, 1, 2, 3...)
- Link function: \log

Response \sim Distribution(Mean Response)

$$Y_i \sim \text{Poisson}(\lambda_i)$$

$$\log(\lambda_i) = a + b x_i$$

$$\lambda_i = e^{a+b x_i}$$

Back to survival of Titanic passengers

How many passengers traveled in each class?

```
tapply(titanic$survived, titanic$class, length)
```

```
#>   crew  first second  third
```

```
#>   885    325    285    706
```

Back to survival of Titanic passengers

How many survived?

```
tapply(titanic$survived, titanic$class, sum)
```

```
#>   crew  first second  third
```

```
#>   212   203   118   178
```

Back to survival of Titanic passengers

What proportion survived in each class?

```
as.numeric(tapply(titanic$survived, titanic$class, mean))  
#> [1] 0.2395480 0.6246154 0.4140351 0.2521246
```

Back to survival of Titanic passengers (package dplyr)

Arrange passenger survival according to class

```
library(dplyr)
summarise(group_by(titanic, class, survived), count = n())
```

Back to survival of Titanic passengers (package dplyr)

Same manipulation using the pipe operator %>%


```
titanic %>%  
  group_by(class, survived) %>%  
  summarise(count = n())
```

Back to survival of Titanic passengers (package dplyr)

Arrange passenger survival according to class

```
#>   class survived    n
#> 1  crew         0 673
#> 2 third         0 528
#> 3  crew         1 212
#> 4 first         1 203
#> 5 third         1 178
#> 6 second        0 167
#> 7 first         0 122
#> 8 second        1 118
```

Or graphically...



Fitting GLMs in R: glm function

```
titanic.glm <- glm(survived ~ class, data = titanic, family = binomial)
#> # A tibble: 4 x 5
#>   term          estimate std.error statistic  p.value
#>   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
#> 1 (Intercept)  -1.16      0.0788    -14.7  1.05e-48
#> 2 classfirst    1.66      0.139     12.0  4.97e-33
#> 3 classecond    0.808     0.144      5.62  1.91e- 8
#> 4 classthird   0.0678    0.117      0.579 5.62e- 1
```

These estimates are on the logit scale!

Interpreting logistic regression outputs

Parameter estimates on the logit scale:

```
#> (Intercept)  classfirst classecond  classtthird  
#> -1.15515905  1.66434399  0.80784987  0.06784632
```

We need to back-transform using the inverse logit function:

```
plogis(coef(titanic.glm)[1]) # crew survival probability  
#> (Intercept)  
#> 0.239548
```

Looking at the data, the proportion of crew who survived is:

```
sum(titanic$survived[titanic$class == "crew"]) /  
  nrow(titanic[titanic$class == "crew", ])  
#> [1] 0.239548
```

Probability of survival for 1st class passengers?

Needs to add intercept (baseline) to the parameter estimate:

```
plogis(coef(titanic.glm)[1] + coef(titanic.glm)[2])  
#> (Intercept)  
#> 0.6246154
```

Again this value matches the data:

```
sum(titanic$survived[titanic$class == "first"]) /  
  nrow(titanic[titanic$class == "first", ])  
#> [1] 0.6246154
```

Generalized linear *mixed* models (or GLMMs)

What are random effects?

- **Mixed models** include both fixed and random effects.
- Random effects are statistical parameters that attempt to **explain noise caused by sub-populations** of the population you are trying to model.
- A random-effect model assumes that the dataset being analysed consists of **a hierarchy of different populations** whose differences relate to that hierarchy.
- Measurement that come **in groups**. E.g. classrooms within schools, chapters within books, populations within species, frogs within ponds.
- Sex or age are not clusters; if we were to sample again, we would take the same levels, e.g. male/female and young/old.

Why do we need random effects?

- Model the clustering itself.
- Interested in variance components (environmental vs. genetic variance)
- Control for bias due to pseudoreplication (time, space, individual)

$$Y_{ij} \sim \text{Distribution}(\text{Mean Response}_{ij})$$

$$\text{Mean Response}_{ij} = \beta_{0j} + \beta_1 x_{i1} + \dots + \beta_P x_{iP}$$

$$\beta_{0j} \sim \text{Normal}(\mu_{group}, \sigma_{group}^2)$$

Model fitting in R


- Linear Mixed Models (LMMs) and Generalized Linear Mixed Models (GLMMs)

```
library(lme4)
my_LMM <- lmer(y ~ x + (1 | group)) # LMM
my_GLMM <- glmer(y ~ x + (1 | group), family = Distribution) # GLMM
```

- x is the **fixed** factor(s)
- 1 is the **random** factor(s), here the intercept
- group is for the **grouping** variable

GLMM with Poisson response

From Jason Matthiopoulos' book



Longitudinal study on coral reef

A survey of a coral reef uses 10 predefined linear transects covered by divers once every week. The response variable of interest is the abundance of a particular species of anemone as a function of water temperature. Counts of anemones are recorded at 20 regular line segments along the transect. The following piece of code will generate a data set with realistic properties according to the above design. Make sure you understand what it is doing. You might want to explain the script to the colleague next to you. Also, to try and make sense of the code of others, it is always good to plot and/or run small sections of the code.

Data generation

```
transects <- 10
data <- NULL
for (tr in 1:transects){
  # random effect (intercept)
  ref <- rnorm(1,0,.5)
  # water temperature gradient
  t <- runif(1, 18,22) + runif(1,-.2,0.2)*1:20
  # Anemone gradient (expected response)
  ans <- exp(ref -14 + 1.8 * t - 0.045 * t^2)
  # actual counts on 20 segments of the current transect
  an <- rpois(20, ans)
  data <- rbind(data, cbind(rep(tr, 20), t, an))
}
```

- Generate a data set using the anemone code.
- Using NIMBLE and brms, fit a GLMM with quadratic effect of temperature and a random intercept.
- Fit the same model to the same data in a Frequentist framework using function `lme4::glmer()`.
- Compare the estimates.


Make sense of the code

- Always difficult to make sense of the code of others.
- Good to plot and/or run small sections of the code.

Make sense of the code

```
# random effect (intercept)  
ref <- rnorm(1,0,.5)  
# water temperature gradient  
t <- runif(1, 18,22) + runif(1,-.2,0.2)*1:20  
  
plot(t,type='l')
```


Make sense of the code



Make sense of the code

```
# Anemone gradient (expected response)  
ans <- exp(ref -14 + 1.8 * t - 0.045 * t^2)  
plot(t,log(ans),type='l')
```


Make sense of the code



Make sense of the code

```
data <- data.frame(Transect = data[,1],  
                  Temperature = data[,2],  
                  Anemones = data[,3])  
  
plot(data$Temperature, data$Anemones)
```

Make sense of the code



Write down model

$\text{Count}_i \sim \text{Poisson}(\lambda_i)$ [likelihood]

$\log(\lambda_i) = a_{\text{TRANSECT}[i]} + b_1 \text{temp}_i + b_2 \text{temp}_i^2$ [linear model]

$a_j \sim \text{Normal}(\bar{a}, \sigma)$ [prior for varying intercepts]

$\bar{a} \sim \text{Normal}(0, 1.5)$ [prior for population mean]

$\sigma \sim \text{Uniform}(0, 10)$ [prior for standard deviation]

$b_1, b_2 \sim \text{Normal}(0, 1.5)$ [prior for slopes]

NIMBLE

Load package

```
library(nimble)
```

Standardize temperature covariate

```
boo <- data$Temperature
data$Temp <- (boo - mean(boo)) / sd(boo)
head(data)
```

#>	Transect	Temperature	Anemones	Temp
#> 1	1	19.88386	28	-0.09912526
#> 2	1	19.68686	29	-0.19824493
#> 3	1	19.48986	24	-0.29736459
#> 4	1	19.29286	30	-0.39648425
#> 5	1	19.09586	39	-0.49560392
#> 6	1	18.89886	25	-0.59472358

Model

```
model <- nimbleCode({
  for (i in 1:n){
    count[i] ~ dpois(lambda[i])
    log(lambda[i]) <- intercept[transect[i]] + slope[1] * x[i] + slope[2] * pow(x[i],2)
  }
  for (j in 1:nbtransects){
    intercept[j] ~ dnorm(mu.a, sd = sigma.a)
  }
  mu.a ~ dnorm(0, sd = 1.5)
  sigma.a ~ dunif(0, 10)
  slope[1] ~ dnorm(0, sd = 1.5)
  slope[2] ~ dnorm(0, sd = 1.5)
})
```

```
my.constants <- list(n = nrow(data),  
                    nbtransects = transects)  
my.data <- list(x = data$Temp,  
               count = data$Anemones,  
               transect = data$Transect)
```

Initial values

```
init1 <- list(intercept = rnorm(transects), slope = rnorm(2),  
              mu.a = rnorm(1), sigma.a = runif(1))  
init2 <- list(intercept = rnorm(transects), slope = rnorm(2),  
              mu.a = rnorm(1), sigma.a = runif(1))  
initial.values <- list(init1, init2)
```

Parameters to monitor (and save)

```
parameters.to.save <- c("slope", "mu.a", "sigma.a")
```

Set up MCMC details

```
n.iter <- 5000  
n.burnin <- 1000  
n.chains <- 2
```

Run model

```
mcmc.output <- nimbleMCMC(code = model,  
                           data = my.data,  
                           constants = my.constants,  
                           inits = initial.values,  
                           monitors = parameters.to.save,  
                           niter = n.iter,  
                           nburnin = n.burnin,  
                           nchains = n.chains)
```

Post-process results

```
library(MCMCvis)
MCMCsummary(object = mcmc.output, round = 2)
#>           mean    sd  2.5%   50% 97.5% Rhat n.eff
#> mu.a         3.61 0.27  3.04   3.63  4.10 1.00  4485
#> sigma.a      0.79 0.23  0.48   0.74  1.36 1.00   998
#> slope[1]    -0.06 0.03 -0.11  -0.06  0.00 1.02   233
#> slope[2]    -0.19 0.02 -0.23  -0.19 -0.15 1.01   302
```

Get regression coefficients

Convert regression coefficients from scaled to non-scaled and compare to values used to generate data (from <https://stats.stackexchange.com/questions/361995/how-to-convert-coefficients-from-quadratic-function-from-scaled-to-not-scaled-co>)

Pull two chains together:

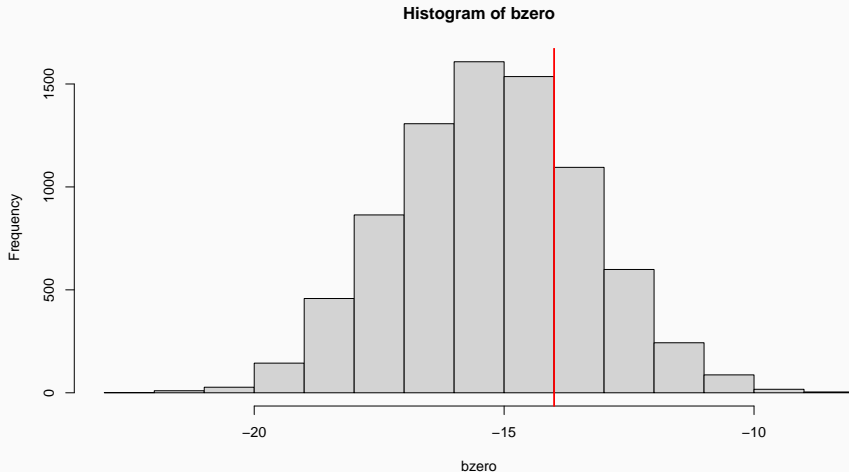
```
samples <- rbind(mcmc.output$chain1,mcmc.output$chain2)
```


Get regression coefficients

```
sbzero <- samples[, 'mu.a']  
sbun <- samples[, 'slope[1]']  
sbdeux <- samples[, 'slope[2]']  
  
mu <- mean(boo)  
sg <- sd(boo)
```

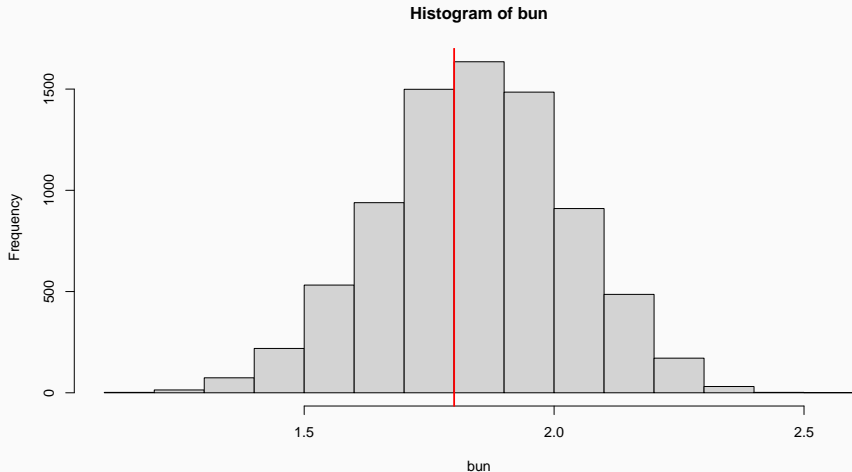
Get regression coefficients

```
bzero <- sbzero - sbun*mu/sg + sbdeux*mu^2/(sg^2)
hist(bzero)
abline(v = -14, col = "red", lwd = 2)
```




Get regression coefficients

```
bun <- sbdeux/sg - 2 * sbdeux * mu / (sg^2)
hist(bun)
abline(v = 1.8, col = "red", lwd = 2)
```



Get regression coefficients

```
bdeux <- sbdeux/(sg^2)
hist(bdeux)
abline(v = - 0.045, col = "red", lwd = 2)
```




Frequentist approach

lme4 fit

```
library(lme4)
fit_lme4 <- glmer(Anemones ~ Temp + I(Temp^2) + (1 | Transect), data = data, family = poisson)
fit_lme4
#> Generalized linear mixed model fit by maximum likelihood (Laplace
#> Approximation) [glmerMod]
#> Family: poisson ( log )
#> Formula: Anemones ~ Temp + I(Temp^2) + (1 | Transect)
#> Data: data
#>      AIC      BIC    logLik deviance df.resid
#> 1357.6583 1370.8516 -674.8292 1349.6583      196
#> Random effects:
#> Groups   Name      Std.Dev.
#> Transect (Intercept) 0.639
#> Number of obs: 200, groups: Transect, 10
#> Fixed Effects:
#> (Intercept)      Temp    I(Temp^2)
#>    3.72064    -0.05814    -0.19002
```

Visualize

```
visreg::visreg(fit_lme4, xvar = 'Temp')
```



brms

Implementation

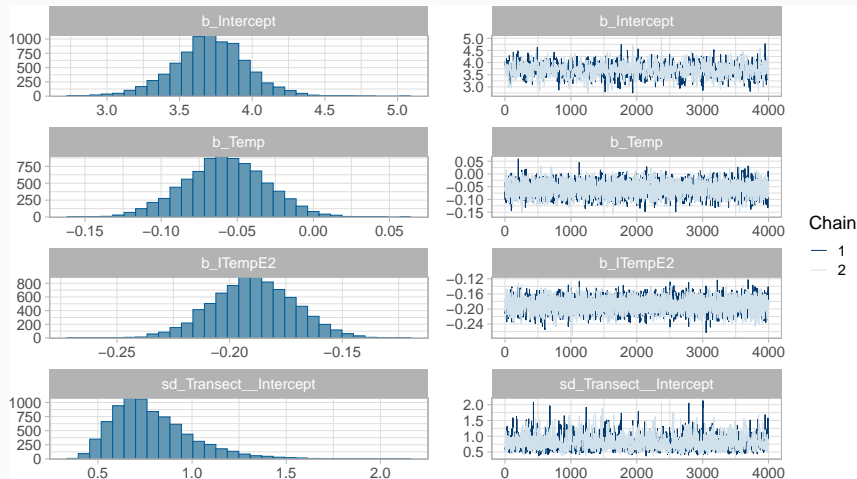
```
library(brms)
bayes.brms <- brm(Anemones ~ Temp + I(Temp^2) + (1 | Transect),
  data = data,
  family = poisson("log"),
  chains = 2, # nb of chains
  iter = 5000, # nb of iterations, including burnin
  warmup = 1000, # burnin
  thin = 1)
```

Numerical summaries

```
summary(bayes.brms)
#> Family: poisson
#> Links: mu = log
#> Formula: Anemones ~ Temp + I(Temp^2) + (1 | Transect)
#> Data: data (Number of observations: 200)
#> Draws: 2 chains, each with iter = 5000; warmup = 1000; thin = 1;
#> total post-warmup draws = 8000
#>
#> Multilevel Hyperparameters:
#> ~Transect (Number of levels: 10)
#> Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
#> sd(Intercept) 0.79 0.22 0.48 1.31 1.00 1511 2441
#>
#> Regression Coefficients:
#> Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
#> Intercept 3.71 0.25 3.18 4.21 1.00 1160 1693
#> Temp -0.06 0.03 -0.11 -0.01 1.00 4284 4289
#> ITempE2 -0.19 0.02 -0.23 -0.15 1.00 4876 4363
#>
#> Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
#> and Tail_ESS are effective sample size measures, and Rhat is the potential
#> scale reduction factor on split chains (at convergence, Rhat = 1).
```


Convergence

```
plot(bayes.brms)
```



Quality of fit assessment

```
pp_check(bayes.brms, ndraws = 100, type = 'ecdf_overlay')
```



Model comparison

- What if we'd like to test the effect of temperature using WAIC?
- We fit a model with no effect of temperature:

```
bayes.brms2 <- brm(Anemones ~ 1 + (1 | Transect),  
  data = data,  
  family = poisson("log"),  
  chains = 2, # nb of chains  
  iter = 5000, # nb of iterations, including burnin  
  warmup = 1000, # burnin  
  thin = 1)
```

Model comparison

Then we compare both models, by ranking them with their WAIC:

```
waic1 <- waic(bayes.brms) # waic model w/ temperature
waic2 <- waic(bayes.brms2) # waic model wo/ temperature
data.frame(with_temp = waic1$waic,
            without_temp = waic2$waic)
#>   with_temp without_temp
#> 1  1308.948    1414.755
```

GLMM with binomial response

Received: 5 July 2021 | Revised: 10 August 2021 | Accepted: 30 September 2021

DOI: 10.1111/ele.13915

LETTER

ECOLOGY LETTERS  WILEY

Sampling bias exaggerates a textbook example of a trophic cascade



Elaine M. Brice¹  | Eric J. Larsen² | Daniel R. MacNulty¹ 

¹Department of Wildland Resources and Ecology Center, Utah State University, Logan, Utah, USA

²Department of Geography and Geology, University of Wisconsin – Stevens Point, Stevens Point, Wisconsin, USA

Correspondence

Elaine M. Brice, Department of Wildland Resources and Ecology Center, Utah State University, Logan, UT 84322, USA.
Email: elbrice92@gmail.com

Funding information

National Science Foundation, Grant/Award Number: DGE-1633756; Yellowstone National Park; University of Wyoming-National Park Service Small Grant Program, Grant/Award Number: 1003867-USU; Utah State University

Editor: Tim Coulson

Abstract

Understanding trophic cascades in terrestrial wildlife communities is a major challenge because these systems are difficult to sample properly. We show how a tradition of non-random sampling has confounded this understanding in a textbook system (Yellowstone National Park) where carnivore [*Canis lupus* (wolf)] recovery is associated with a trophic cascade involving changes in herbivore [*Cervus canadensis* (elk)] behaviour and density that promote plant regeneration. Long-term data indicate a practice of sampling only the tallest young plants overestimated regeneration of overstory aspen (*Populus tremuloides*) by a factor of 4–7 compared to random sampling because it favoured plants taller than the preferred browsing height of elk and overlooked non-regenerating aspen stands. Random sampling described a trophic cascade, but it was weaker than the one that non-random sampling described. Our findings highlight the critical importance of basic sampling principles (e.g. randomisation) for achieving an accurate understanding of trophic cascades in terrestrial wildlife systems.

KEYWORDS

aspen, carnivore, elk, non-random sampling, predator indirect effects, preferred browsing height, sampling bias, trophic cascade, ungulate, wolf

Read in data

```
rawdat <- read_csv("dat/Aspen_Data.csv")
head(rawdat)
#> # A tibble: 6 x 6
#>   Plot Year Tree Browse Height Type
#>   <dbl> <dbl> <dbl>  <dbl>  <dbl> <chr>
#> 1     1     1 2008     1      1   115 5T
#> 2     1     1 2008     2      1   130 5T
#> 3     1     1 2008     3      1   113 5T
#> 4     1     1 2008     4      1   110 5T
#> 5     1     1 2008     5      0   132 5T
#> 6     1     1 2009     1      1   200 5T
```

Data collection (1)

We measured browsing and height of young aspen in 113 plots distributed randomly across the study area. Each plot was a 1 x 20 m belt transect located randomly within an aspen stand that was itself randomly selected from an inventory of stands with respect to high and low wolf-use areas. The inventory was a list of 992 grid cells (240 × 360 m) that contained at least one stand. A 'stand' was a group of tree-size aspen (more than 10 cm diameter at breast height) in which each tree was less than 30 m from every other tree. One hundred and thirteen grid cells were randomly selected from the inventory, one stand was randomly selected from each cell, and one plot was randomly established in each stand.

Data collection (2)

We measured aspen at the end of the growing season (late July–September), focusing on plants less than 600 cm tall and more than 1 year old, which we termed ‘young aspen’. For each stand, we measured every young aspen within a plot (‘random stems’) and each of the five tallest young aspen within the stand (‘5T stems’). For all young aspen, we measured browsing status (browsed or unbrowsed) and height of the leader (tallest) stem. A leader was ‘browsed’ if its growth from the previous growing season had been eaten, which we identified by a sharp, pruned edge at the base of the current year’s growth. Most plots were measured nearly every year since 1999 and our analysis focused on data from 10 years (2007–2014, 2016–2017) in which sampled stands included measurements of random and 5T stems.

We combined measurements of 5T and random stems into one data set of all stems ($N = 18,623$) across all years ($N = 10$ years).

- Read in the data and get familiar with it.
- Using NIMBLE and brms, fit a GLMM with 'Browsed' as the response variable. Consider year and type of stems (random or 5T stems) as explanatory variables. Use stand identity as a random effect on the intercept.
- Fit the same model to the same data in a Frequentist framework using function `lme4::glmer()`.
- Compare the estimates.

Read in data

```
str(rawdat)
#> spc_tbl_ [18,792 x 6] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
#> $ Plot   : num [1:18792] 1 1 1 1 1 1 1 1 1 1 ...
#> $ Year   : num [1:18792] 2008 2008 2008 2008 2008 ...
#> $ Tree   : num [1:18792] 1 2 3 4 5 1 2 3 4 5 ...
#> $ Browse: num [1:18792] 1 1 1 1 0 1 0 0 0 0 ...
#> $ Height: num [1:18792] 115 130 113 110 132 200 208 218 240 260 ...
#> $ Type   : chr [1:18792] "5T" "5T" "5T" "5T" ...
#> - attr(*, "spec")=
#> .. cols(
#> ..   Plot = col_double(),
#> ..   Year = col_double(),
#> ..   Tree = col_double(),
#> ..   Browse = col_double(),
#> ..   Height = col_double(),
#> ..   Type = col_character()
#> .. )
#> - attr(*, "problems")=<externalptr>
```

Metadata

- Plot: individual identifier for each of 113 plots distributed randomly across the study area. Each plot was a 1×20 m belt transect located randomly within an aspen stand
- Year: year in which aspen was sampled
- Tree: individual identifier for each stem within a plot
- Browse: denotes the browsing status (browsed = 1, unbrowsed = 0) of the leader (tallest) stem. A leader was 'browsed' if its growth from the previous growing season had been eaten
- Height: height (cm) of the leader stem of each individual aspen
- Type: sampling method. Every young aspen within a plot is a "random" stem, and each of the five tallest young aspen within the stand is a "5T" stem.

Format data

```
dat <- rawdat %>%  
  transmute(y = Browse,  
            type = if_else(Type == "5T", 1, 0), # numeric  
            year = as.integer(Year) - 2006,      # year starts at 1  
            id = as_factor(Plot))                # id is a factor
```


Write down model

$\text{Browsed}_i \sim \text{Bernoulli}(p_i)$	[likelihood]
$\text{logit}(p_i) = a_{\text{ID}[i]} + b_1 \text{ year}_i + b_2 \text{ type}_i$	[linear model]
$a_j \sim \text{Normal}(\bar{a}, \sigma)$	[prior for varying intercepts]
$\bar{a} \sim \text{Normal}(0, 1.5)$	[prior for population mean]
$\sigma \sim \text{Uniform}(0, 10)$	[prior for standard deviation]
$b_1, b_2 \sim \text{Normal}(0, 1.5)$	[prior for slopes]

NIMBLE

Load package

```
library(nimble)
```

Model

```
model <- nimbleCode({
  for (i in 1:n){
    y[i] ~ dbern(p[i])
    logit(p[i]) <- a[id[i]] + beta[1] * type[i] + beta[2] * year[i] + beta[3] * year[i] * type[i]
  }
  for (j in 1:nlevels){
    a[j] ~ dnorm(mu.a, sd = sigma.a)
  }
  for (k in 1:3){
    beta[k] ~ dnorm(0, sd = 1.5)
  }
  mu.a ~ dnorm(0, sd = 1.5)
  sigma.a ~ dunif(0,10)
})
```

```
my.constants <- list(id = as.numeric(as_factor(dat$id)),  
                    n = length(dat$y),  
                    nlevels = length(levels(as_factor(dat$id))))  
my.data <- list(y = dat$y,  
              type = dat$type,  
              year = dat$year)
```

Initial values

```
init1 <- list(mu.a = -0.5, sigma.a = 0.1, beta = rnorm(3))  
init2 <- list(mu.a = 0.5, sigma.a = 0.5, beta = rnorm(3))  
initial.values <- list(init1, init2)
```

Parameters to monitor (and save)

```
parameters.to.save <- c("mu.a", "sigma.a", "beta")
```

Set up MCMC details

```
n.iter <- 5000  
n.burnin <- 1000  
n.chains <- 2
```


Run model

```
mcmc.output <- nimbleMCMC(code = model,  
                           data = my.data,  
                           constants = my.constants,  
                           inits = initial.values,  
                           monitors = parameters.to.save,  
                           niter = n.iter,  
                           nburnin = n.burnin,  
                           nchains = n.chains)
```

Post-process results

```
library(MCMCvis)
MCMCsummary(object = mcmc.output, round = 2)
#>           mean    sd  2.5%   50% 97.5% Rhat n.eff
#> beta[1] -0.43 0.09 -0.61 -0.43 -0.25 1.00   186
#> beta[2] -0.14 0.01 -0.15 -0.14 -0.12 1.02   130
#> beta[3] -0.08 0.01 -0.11 -0.08 -0.06 1.00   179
#> mu.a      1.57 0.13  1.31  1.57  1.82 1.00 1494
#> sigma.a   1.26 0.10  1.08  1.25  1.48 1.00  980
```

Visualize (1)

```
# pool two chains together
samples <- rbind(mcmc.output$chain1, mcmc.output$chain2)
nsim <- nrow(samples)

# get values sampled in posterior distribution of regression parameters
beta1 <- samples[, 'beta[1]']
beta2 <- samples[, 'beta[2]']
beta3 <- samples[, 'beta[3]']
intercept <- samples[, 'mu.a']

# predict
logitp_5T <- matrix(NA, nrow = nsim, ncol = 11)
logitp_random <- matrix(NA, nrow = nsim, ncol = 11)
for (i in 1:nsim){ # loop over simulations
  for (j in 1:11){ # loop over years
    logitp_5T[i,j] <- intercept[i] + beta1[i] * 1 + beta2[i] * j + beta3[i] * 1 * j
    logitp_random[i,j] <- intercept[i] + beta1[i] * 0 + beta2[i] * j + beta3[i] * 0 * j
  }
}

# back-transform
p_5T <- plogis(logitp_5T)
p_random <- plogis(logitp_random)
```

Visualize (2)


```
# means and credible intervals
p_5T_mean <- apply(p_5T, 2, mean)
p_5T_inf <- apply(p_5T, 2, quantile, probs = 2.5/100)
p_5T_sup <- apply(p_5T, 2, quantile, probs = 97.5/100)

p_random_mean <- apply(p_random, 2, mean)
p_random_inf <- apply(p_random, 2, quantile, probs = 2.5/100)
p_random_sup <- apply(p_random, 2, quantile, probs = 97.5/100)
```

Visualize (3)

```
data.frame(year = c(1:11, 1:11),
            type = factor(c(rep(1, 11), rep(0, 11))), # Convert type to a factor
            p_mean = c(p_5T_mean, p_random_mean),
            p_inf = c(p_5T_inf, p_random_inf),
            p_sup = c(p_5T_sup, p_random_sup)) %>%
ggplot() +
geom_ribbon(aes(x = year,
               ymin = p_inf,
               ymax = p_sup,
               group = type,
               fill = type),
           alpha = 0.3,
           show.legend = F) +
geom_line(aes(x = year,
              y = p_mean,
              group = type,
              color = type),
          size=0.7) +
scale_color_manual(values = c("1" = "#CC2529", "0" = "#396AB1")) +
scale_fill_manual(values = c("1" = "#CC2529", "0" = "#396AB1")) +
xlab("Year")+
ylab("Browse probability")+
theme_light()
```

Visualize (3)




Frequentist approach

lme4 fit

```
library(lme4)
fit_lme4 <- glmer(y ~ type * year + (1|id), data = dat, family = "binomial")
fit_lme4
#> Generalized linear mixed model fit by maximum likelihood (Laplace
#> Approximation) [glmerMod]
#> Family: binomial (logit)
#> Formula: y ~ type * year + (1 | id)
#> Data: dat
#>      AIC      BIC    logLik deviance df.resid
#> 21428.86 21468.02 -10709.43  21418.86    18618
#> Random effects:
#> Groups Name      Std.Dev.
#> id      (Intercept) 1.233
#> Number of obs: 18623, groups: id, 113
#> Fixed Effects:
#> (Intercept)      type      year  type:year
#>    1.57173    -0.41777    -0.13750    -0.08506
```


Visualize

```
visreg::visreg(fit = fit_lme4, xvar = "year",  
               by = "type", overlay = TRUE, scale = "response")
```



brms

Implementation

```
library(brms)
bayes.brms <- brm(y ~ type * year + (1|id),
  data = dat,
  family = bernoulli("logit"),
  chains = 2, # nb of chains
  iter = 5000, # nb of iterations, including burnin
  warmup = 1000, # burnin
  thin = 1)
```

Numerical summaries

```
summary(bayes.brms)
#> Family: bernoulli
#> Links: mu = logit
#> Formula: y ~ type * year + (1 | id)
#> Data: dat (Number of observations: 18623)
#> Draws: 2 chains, each with iter = 5000; warmup = 1000; thin = 1;
#> total post-warmup draws = 8000
#>
#> Multilevel Hyperparameters:
#> ~id (Number of levels: 113)
#>      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
#> sd(Intercept)    1.25     0.10    1.08    1.46 1.00     947    1197
#>
#> Regression Coefficients:
#>      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
#> Intercept      1.57     0.13    1.31    1.83 1.00     404     810
#> type          -0.42     0.09   -0.59   -0.24 1.00    5716    5214
#> year          -0.14     0.01   -0.15   -0.12 1.00    8197    5558
#> type:year     -0.09     0.01   -0.11   -0.06 1.00    5668    5123
#>
#> Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS
#> and Tail_ESS are effective sample size measures, and Rhat is the potential
#> scale reduction factor on split chains (at convergence, Rhat = 1).
```

Convergence

```
plot(bayes.brms)
```

